


## Module 1: Advanced Docker — Detailed Notes for Participants

### 1. Multi-Stage Builds & Image Optimization



```
# Base image Stage 1
FROM ubuntu:16.04 as stage1

RUN apt-get update
RUN apt-get -y install make curl
RUN curl http://xyz.com/abc.tar.gz -O
RUN tar zxf abc.tar.gz && cd abc
RUN make DESTDIR=/tmp install

# Stage 2
FROM alpine:3.10

COPY --from=stage1 /tmp /abc

ENTRYPOINT ["/abc/app"]
```

## 1.1 Why Multi-Stage Builds?

Multi-stage builds help you:

- **Reduce image size**
- **Speed up deployments**
- **Improve security** by excluding unnecessary dev tools
- **Separate build & runtime environments**

Traditional Dockerfiles produced large images because they included compilers, dependencies, and runtime components in a single image.

## 1.2 How Multi-Stage Builds Work

You use multiple FROM statements in the same Dockerfile—each stage creates intermediate images.

Example:

```
# Stage 1: Builder
```

```
FROM golang:1.21 AS builder
```

```
WORKDIR /app
COPY .
RUN go build -o app .

# Stage 2: Runtime
FROM alpine:3.18
WORKDIR /app
COPY --from=builder /app/app .
CMD ["./app"]
```

### 1.3 Benefits

- Final image contains **only what is required** to run the application.
- No need to install compilers in the runtime image.
- Helps in **CI/CD pipelines** where speed and security matter.

### 1.4 Best Practices

- Use **Alpine** or other lightweight base images.
  - Use `.dockerignore` to avoid copying unnecessary files.
  - Pin versions: `node:20-alpine` instead of `node:latest`.
- 

## 2. Docker Compose for Production

Docker Compose is not only for local development. With proper configurations, it works well for **small/medium-scale production** setups.

### 2.1 YAML Structure Refresher

A Compose file (`docker-compose.yml`) describes:

- Services
- Networks
- Volumes
- Dependencies
- Environment variables
- Build settings

## 2.2 Example Production-Ready Compose

```
version: "3.9"
```

```
services:
```

```
  web:
```

```
    image: myapp:1.0.0
```

```
    ports:
```

```
      - "80:80"
```

```
    environment:
```

```
      - APP_ENV=prod
```

```
    depends_on:
```

```
      - db
```

```
  restart: unless-stopped
```

```
networks:
```

```
  - app-net
```

```
db:
```

```
  image: postgres:15
```

```
  volumes:
```

```
    - dbdata:/var/lib/postgresql/data
```

```
  environment:
```

```
    POSTGRES_USER: admin
```

```
    POSTGRES_PASSWORD: StrongP@ss
```

```
  networks:
```

```
    - app-net
```

```
volumes:
```

```
  dbdata:
```

```
networks:
```

```
  app-net:
```

### 2.3 Production Tips

- Always use **restart policies** (e.g., on-failure, unless-stopped)
- Store secrets using **Docker secrets**, not environment variables.
- Use **named volumes** to persist data.
- Configure **healthcheck** to restart unhealthy containers.

Example:

healthcheck:

```
test: ["CMD", "curl", "-f", "http://localhost/health"]  
interval: 30s  
timeout: 10s  
retries: 3
```

### 2.4 Scaling


Use:

```
docker compose up --scale web=3 -d
```

This launches **horizontal scaling**, useful for load balancing.

---

## 3. Docker Volumes, Networking & Secrets Management



### **3.1 Docker Volumes**

Volumes allow data persistence.

#### **Types of volumes**

1. **Named volumes** (Recommended)
2. **Anonymous volumes**
3. **Bind mounts** (For development)

Example:

```
docker volume create appdata
```

```
docker run -v appdata:/data alpine
```

#### **When to use volumes**

- Databases
- Config files
- Logs that must persist across restarts

---

### **3.2 Docker Networking**

Docker provides 3 major network types:

<b>Network Type</b>	<b>Description</b>	<b>Use Case</b>
<b>bridge</b>	Default network	Multi-container apps
<b>host</b>	Shares host network	High performance apps
<b>overlay</b>	Multi-host networking	Docker Swarm/K8s

#### **Inspect networks**

```
docker network ls
```

```
docker network inspect bridge
```

#### **Create a custom network**

```
docker network create mynet
```

```
docker run --network=mynet nginx
```

Benefits:

- Better **isolation**
- DNS-based **service discovery**
- Container-to-container communication

---

### 3.3 Docker Secrets Management

Secrets include:

- API keys
- DB passwords
- TLS certificates
- OAuth tokens

#### Why not store secrets in environment variables?

- They appear in container metadata.
- Anyone with Docker access can view them.
- They may be logged accidentally.

#### Secure method (Docker Swarm required)

```
echo "SecretP@ss" | docker secret create db_pass -  
docker service create --name myapp --secret db_pass nginx
```

Inside container, secrets are mounted at:


```
/run/secrets/db_pass
```



#### Best Practices

- Rotate secrets regularly.
- Never commit secrets to Git.
- Use .env + secrets manager (AWS KMS, Vault, Docker Swarm).

---

## 4. Scanning & Securing Docker Images





6

#### 4.1 Why Scan Docker Images?

Containers often include:

- CVEs (Common Vulnerabilities & Exposures)
- Outdated libraries
- Misconfigured permissions
- Hardcoded secrets

#### 4.2 Popular Scanning Tools

Tool	Purpose
Trivy	Vulnerabilities + misconfig + secrets
Grype	Fast CVE scanner
Clair	Registry-level scanning

Tool	Purpose
Docker Scout	Built into Docker Hub
Anchore	CI/CD scanning

---

#### 4.3 Example: Scan with Trivy

Install:

```
sudo apt install trivy
```

Scan an image:

```
trivy image nginx:latest
```

Scan file system for secrets:

```
trivy fs --scanners secret .
```

#### 4.4 Security Best Practices

- Use **official images** whenever possible.
- Always pin versions: python:3.10-slim.
- Never run containers as **root** (use USER 1000).
- Use **read-only root filesystem**:

read\_only: true

- Enable **seccomp, AppArmor, SELinux**.
- Limit container capabilities:

cap\_drop:

- ALL

- Use **image signing** with cosign.

#### 4.5 CI/CD Integration

Add scanning steps in pipeline:

- GitHub Actions
- GitLab CI
- Jenkins
- Azure DevOps

Example (Jenkins):

```
trivy image --exit-code 1 myapp:1.0.0
```

If vulnerabilities found, build fails.

---

### ✓ Module 1 Summary

Topic	Quick Recap
Multi-Stage Builds	Reduce image size, split build/runtime
Docker Compose Production	Scaling, healthchecks, restart policies
Volumes	Persistent storage
Networking	Custom networks, DNS-based discovery
Secrets	Secure storage using Swarm/KMS/Vault
Scanning	Use Trivy/Grype to identify CVEs, misconfigs, secrets